

Звіт

Виконала студентка групи ПП-41/1

Терещук Анна

Перед початком виконання основних завдань проводимо дослідницький аналіз даних:

```
In [37]: data_frame = pd.read_csv("Automobile_data.csv")
display(data_frame)
data_frame[data_frame["normalized-losses"] == "?"]
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	1
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	1
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	1
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.4	10.0	1
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.4	8.0	1
...
200	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	1
201	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	1
202	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	1
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.4	23.0	1
204	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	1

205 rows × 26 columns

```
In [13]: data_frame.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column              Non-Null Count  Dtype
---  -
0   symboling            205 non-null    int64
1   normalized-losses    205 non-null    object
2   make                 205 non-null    object
3   fuel-type            205 non-null    object
4   aspiration            205 non-null    object
5   num-of-doors         205 non-null    object
6   body-style           205 non-null    object
7   drive-wheels         205 non-null    object
8   engine-location      205 non-null    object
9   wheel-base          205 non-null    float64
10  length               205 non-null    float64
11  width                205 non-null    float64
12  height               205 non-null    float64
13  curb-weight          205 non-null    int64
14  engine-type          205 non-null    object
15  num-of-cylinders     205 non-null    object
16  engine-size          205 non-null    int64
17  fuel-system          205 non-null    object
18  bore                 205 non-null    object
19  stroke               205 non-null    object
20  compression-ratio    205 non-null    float64
21  horsepower           205 non-null    object
22  peak-rpm             205 non-null    object
23  city-mpg             205 non-null    int64
24  highway-mpg          205 non-null    int64
25  price                205 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB
```

```
In [14]: data_frame.isnull().sum()
```

```
Out[14]: symboling            0
normalized-losses          0
make                        0
fuel-type                   0
aspiration                  0
num-of-doors                0
body-style                  0
drive-wheels                0
engine-location             0
wheel-base                 0
length                      0
width                       0
height                      0
curb-weight                 0
engine-type                 0
num-of-cylinders            0
engine-size                 0
fuel-system                 0
bore                        0
stroke                      0
compression-ratio           0
horsepower                  0
peak-rpm                    0
city-mpg                    0
highway-mpg                 0
price                       0
dtype: int64
```

```
In [15]: data_frame.duplicated().sum()
```

```
Out[15]: 0
```

```
In [17]: data_frame
```

```
Out[17]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	1
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	1
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	1
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.4	10.0	1
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.4	8.0	1
...
200	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	1
201	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	1
202	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	1
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.4	23.0	1
204	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	1

205 rows × 26 columns

Використовуємо методи для оцінки значущості факторів.

```
In [23]: num_val = transformed_data.select_dtypes(exclude='object')
obj_val = transformed_data.select_dtypes(exclude='int64')
obj_val = transformed_data.select_dtypes(exclude='float64')

varianceThreshold = VarianceThreshold(threshold=(0.5))
varianceThreshold.fit(num_val)
selected_data = varianceThreshold.transform(num_val)
selected_data

Out[23]: array([[1.6400e+02, 0.0000e+00, 3.0000e+00, ..., 2.4000e+01, 3.0000e+01,
1.3950e+04],
[1.6400e+02, 0.0000e+00, 3.0000e+00, ..., 1.8000e+01, 2.2000e+01,
1.7450e+04],
[1.5800e+02, 0.0000e+00, 3.0000e+00, ..., 1.9000e+01, 2.5000e+01,
1.7710e+04],
...,
[9.5000e+01, 1.7000e+01, 3.0000e+00, ..., 1.8000e+01, 2.3000e+01,
2.1485e+04],
[9.5000e+01, 1.7000e+01, 3.0000e+00, ..., 2.6000e+01, 2.7000e+01,
2.2470e+04],
[9.5000e+01, 1.7000e+01, 3.0000e+00, ..., 1.9000e+01, 2.5000e+01,
2.2625e+04]])

In [24]: num_val
Out[24]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
3	0.8	164.0	0	1	0	0	3	1	0	99.8	109	4	3.19	3.40	10.0	102	5500
4	0.8	164.0	0	1	0	0	3	0	0	99.4	136	4	3.19	3.40	8.0	115	5500
6	0.6	158.0	0	1	0	0	3	1	0	105.8	136	4	3.19	3.40	8.5	110	5500
8	0.6	158.0	0	1	1	0	3	1	0	105.8	131	4	3.13	3.40	8.3	140	5500
10	0.8	192.0	1	1	0	1	3	2	0	101.2	108	4	3.50	2.80	8.8	101	5500
...
200	0.2	95.0	17	1	0	0	3	2	0	109.1	141	4	3.78	3.15	9.5	114	5500
201	0.2	95.0	17	1	1	0	3	2	0	109.1	141	4	3.78	3.15	8.7	160	5500
202	0.2	95.0	17	1	0	0	3	2	0	109.1	173	4	3.58	2.87	8.8	134	5500
203	0.2	95.0	17	0	1	0	3	2	0	109.1	145	2	3.01	3.40	23.0	106	4800
204	0.2	95.0	17	1	1	0	3	2	0	109.1	141	4	3.78	3.15	9.5	114	5500

159 rows × 26 columns

```
In [25]: obj_val
Out[25]:
```

	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	curb-weight	engine-type	num-of-cylinders	engine-size	fuel-system	horsepower	peak-rpm	city-mpg	highway-mpg
3	0	1	0	0	3	1	0	2337	2	2	109	4	102	5500	24	30
4	0	1	0	0	3	0	0	2824	2	1	136	4	115	5500	18	22
6	0	1	0	0	3	1	0	2844	2	1	136	4	110	5500	19	25
8	0	1	1	0	3	1	0	3086	2	1	131	4	140	5500	17	20
10	1	1	0	1	3	2	0	2395	2	2	108	4	101	5800	23	29
...
200	17	1	0	0	3	2	0	2952	2	2	141	4	114	5400	23	28
201	17	1	1	0	3	2	0	3049	2	2	141	4	160	5300	19	25
202	17	1	0	0	3	2	0	3012	4	3	173	4	134	5500	18	23
203	17	0	1	0	3	2	0	3217	2	3	145	2	106	4800	26	27
204	17	1	1	0	3	2	0	3062	2	2	141	4	114	5400	19	25

159 rows × 17 columns

Наступний метод обирає k найвпливовіших факторів:

```
In [28]: columns_array = transformed_data.columns
kbest = sklearn.feature_selection.SelectKBest(sklearn.feature_selection.chi2, k=5)
transformed_data_kbest = kbest.fit_transform(num_val.drop(["price"], axis=1),
num_val["price"])
selected_features = kbest.get_feature_names_out(columns_array[:len(columns_array) - 1])
transformed_data_kbest = pd.DataFrame(transformed_data_kbest, columns=selected_features)
transformed_data_kbest


Out[28]:
```

	normalized-losses	curb-weight	engine-size	horsepower	peak-rpm
0	164.0	2337.0	109.0	102.0	5500.0
1	164.0	2824.0	136.0	115.0	5500.0
2	158.0	2844.0	136.0	110.0	5500.0
3	158.0	3086.0	131.0	140.0	5500.0
4	192.0	2395.0	108.0	101.0	5800.0
...
154	95.0	2952.0	141.0	114.0	5400.0
155	95.0	3049.0	141.0	160.0	5300.0
156	95.0	3012.0	173.0	134.0	5500.0
157	95.0	3217.0	145.0	106.0	4800.0
158	95.0	3062.0	141.0	114.0	5400.0

159 rows × 5 columns

Зображуємо важливість кожної метрики за допомогою RandomForestRegressor:

```
In [33]: feature_importances = tree_classifier.feature_importances_  
indices = np.argsort(feature_importances)  
plt.xticks(range(len(indices)), np.array(num_val.drop(["price"], axis=1).columns)[indices])  
plt.barh(range(len(indices)), feature_importances[indices])  
plt.show()
```



На наступному графіку можна побачити загальну кореляцію метрик, що відображає значення із найбільшою кореляцією:

Out[36]: <AxesSubplot: >

